

**Everyone is reading
this leaderboard
wrong.**

The score column isn't the one that matters.

Here's the number nobody mentioned.

Opus 4.7 Max	64.8%	\$11.02/task
Composer 2.5	63.2%	\$0.55/task

1.6 point score difference. 20x cost difference.

Run 100 tasks a day. Here's the bill.

Opus 4.7 Max	\$1,102/day	\$402K/year
Composer 2.5	\$55/day	\$20K/year

\$382,000 extra per year for 1.6% more accuracy.

Same model. Different budget.

Opus 4.7 runs across 5 tiers. Max, Extra High, High, Medium, Low. Same underlying model. More thinking tokens = higher score = higher cost.

Max	64.8%	\$11.02
High	59.4%	\$5.01
Low	48.3%	\$1.87

Max

64.8%

\$11.02

High

59.4%

\$5.01

Low

48.3%

\$1.87

You're not buying a smarter model.

You're buying the same model with more time to think.
You're paying for patience.

The Composer anomaly.

Third place overall. \$0.55 per task. Twenty cents cheaper than Gemini 3.5 Flash, which scores 14 points lower.

Composer 2.5 score

63.2%

Gemini 3.5 Flash score

49.8%

Cost difference

\$1.39

Score per dollar. That's the ratio.

Capital efficiency. Same question finance asks about every trade.

Composer 2.5	115 pts/\$	best
GPT-5.5 Medium	27 pts/\$	good
Opus 4.7 Max	5.9 pts/\$	costly

Most tasks don't need the top tier.

The builders who figure this out early have systems that actually scale. The ones who don't have expensive demos.

FOLLOW FOR MORE FROM INSIDE FINANCE AND AI.

Read the cost column.

That's usually where the actual decision lives.

Full breakdown at bionicbanker.tech